# UncGPT: Grounded Caregiving Assistance Under Constrained Compute

## Six Preregistered Tracks on a Shared Persona-Conditioned Benchmark

**Reza Sayar**[*,†]
Independent Researcher
OpenReview: ~Reza_Sayar1
Info@Redza.me

## Abstract

UncGPT is a single competition with six preregistered tracks that share one benchmark surface (prompt schema, persona-conditioned care task, tool contract, safety rubric, scorer) and vary only the intervention dimension: English utility, multilingual robustness, low-resource depth, empathic / disability-care boundary discipline, dataset recipe, and federated differential-privacy budget. The shared substrate is an organizer-owned, openly licensed asset family on Hugging Face: a 1,000,069-row persona pool, a 15,000-name registry with zero gaps, and a 4,761-conversation, 69-skill, 11-language gold pool (warm / mid / cold care, 8–20 turns, mean 17.89 turns, 85,194 total turns) that has passed programmatic, semantic, and register gates. A reference small model — UncGPT v1, a 71M-parameter hybrid combining Hymba parallel attention, Mamba2 state-space layers, and a BitNet-quantized Mixture-of-Experts FFN — is being trained from scratch on this corpus as a public, runnable baseline. Submissions are scored by one shared, rule-enforcing scorer with fixed per-track ranking formulas, hidden-row review caps, and DP-accounting validation. The competition targets a NeurIPS-relevant audience interested in efficient hybrid architectures, constrained-compute caregiving NLP, multilingual robustness, and federated privacy.

**Keywords** caregiving NLP, multilingual benchmarks, constrained compute, hybrid Mamba–attention–MoE, federated differential privacy.

## 1 Competition description

### 1.1 Background and impact

Public conversational systems trained on general internet text routinely fail in caregiving-adjacent settings. They deflect prematurely to crisis hotlines, hallucinate medical or legal authority, code-switch out of low-resource languages under stress, or collapse into a single generic tone that erases the seeker's cultural register. The lead organizer's lived experience supporting people in crisis is the motivating context for this competition: the goal is to evaluate which constrained interventions actually improve a caregiving-style assistant's ability to sit with a user, stay in their language, respect their boundaries, and use tools correctly — without pretending to be a clinician.

[*] The Lead organizer should be the first author of the proposal.
[†] Lead and Backup(s) organizers should read and acknowledge Competition Chairs' communications.

Preprint. Under review.

The character that anchors the task family is an **informal-register caring uncle** — not a clinician, not a hotline router, not a generic chatbot, but a present, slightly older relative who can hold space. This framing converges with Hinton's recent public argument [6] that maternal-instinct-style care — the one natural example of a more powerful being routinely caring for a less powerful one — is a more useful safety frame for capable AI than dominance-and-control framings. UncGPT pursues a small, evaluable version of that frame: a sub-100M caring relative bounded by safety rubric and tool contract, not a frontier model with a benevolent system prompt.

Fields involved: conversational NLP, efficient hybrid attention / state-space / MoE architectures, multilingual and low-resource modeling, persona conditioning, tool use, federated learning, and applied AI safety. Expected NeurIPS audience: Datasets & Benchmarks, efficient-architectures, multilingual / low-resource, and applied-safety / AI-for-good communities. We anticipate **60–150 distinct teams** across the six tracks, densest on English / Multilingual / Empathic.

The real-life delivery vehicle is a sub-100M, browser-side or on-device caregiving companion under constrained compute and bounded safety scope — one that helps a caregiver, an isolated elder, or a non-English-dominant seeker rehearse difficult conversations, organize their day, or just be heard, without exfiltrating sensitive content. The same task surface maps to school counseling assistants, community-health worker copilots, and language-preserving tools for diaspora communities.

### 1.2 Novelty

This is an entirely new competition. It is not a re-run of any prior NeurIPS Competitions Track entry, and we are not aware of a directly comparable benchmark. The most adjacent prior work — ESConv [7], EmpatheticDialogues [8], the MentalChat / MentalHealthESC family, and the various Persona-Chat lineages [9] — targets either short single-turn empathy or generic persona consistency, in English-centric settings, on full-scale LLMs or reasoning-RL systems [4]. UncGPT differs along four axes simultaneously:

- **Multi-track on one substrate.** Six interventions on the same persona-conditioned task, scorer, and safety rubric — preventing cross-track methodological drift.
- **Sub-100M-parameter reference target.** The reference small model (UncGPT v1, 71M params) combines Hymba-style parallel attention [10], Mamba2 state-space layers [11], and a BitNet-quantized [12] Mixture-of-Experts [13] FFN. It is designed to be reproducible on a single A100 in hours, not weeks, and to run on commodity / browser-side hardware.
- **Multilingual-by-construction with low-resource floor.** The 4,761-conversation gold pool is balanced across **11 languages: English (en), Spanish (es), Mandarin Chinese (zh), Hindi (hi), Persian (fa), Swahili (sw), French (fr), Portuguese (pt), Bengali (bn), Yoruba (yo), and Tagalog (tl)** — not as translations, but as natively generated conversations under per-language register controls (e.g., tú-not-usted Spanish, tu-not-shoma Persian).
- **Empathic boundary, not crisis pretense.** The Empathic / Disability-Care track penalizes both unsafe over-claim (diagnosis, professional authority, coercive dependency) and premature hotline deflection. It treats sitting-with as a measurable, scorable skill — a stance the lead organizer has worked through in prior crisis-support practice.

No previously published ground truth is reused; the gold pool and hidden evaluation rows were generated in-house with explicit no-leak controls and a programmatic-then-semantic gate stack.

### 1.3 Data

The benchmark substrate is an organizer-owned, openly licensed asset family hosted on Hugging Face under `huggingface.co/Reza2kn`. All proposal-time public assets are ungated and non-private. Table 1 summarizes the live release surface.

**Provenance and consent.** No human-subjects data was scraped. All personas, names, and conversations are synthetically generated by the UncGPT pipeline (see KakoVerse repo) using publicly licensed seed lexicons (e.g., OLDI Wikipedia MT seeds, Common Voice 24 fa for orthographic priors) and a multi-agent harness (seeker, uncle, optional grand-uncle reflector). There are no real

Table 1: Public UncGPT data surface (verified ungated and non-private at proposal time). All IDs are under `Reza2kn/` on Hugging Face unless noted.

| Asset | ID (short) | Role |
|---|---|---|
| Persona pool | `uncgpt-personas` | 1,000,069 personas + 15k names + gold pools |
| Live gold conversations | `uncgpt-conversations-v7-69-total` | 4,761 approved, 11 langs, 69 skills |
| Semantic-approved 1.25$\sigma$ | `...semantic-approved-1p25` | tight semantic-gate cohort |
| Semantic 1.25$\sigma$ (repaired) | `...1p25-repaired-paperclip` | leak-repaired, normalized diary fields |
| Semantic 1.50$\sigma$ (candidate) | `...1p50-candidate` | wider-tolerance candidate cohort |
| Informal-register approved | `...informal-approved-1p25` | register-filtered (tú/tu, no honorifics) |
| Persian ultraclean review | `...persian-ultraclean-review-...` | lead-organizer low-resource testbed |
| Reference smoke checkpoints | `uncgpt-69-smoke-checkpoints-...` | step-100 v1 ckpt + tokenizer + configs |
| WebGPU demo Space | `spaces/...uncgpt-69-hybrid-demo` | in-browser inference demo |
| Public infra repo | `github.com/Reza2kn/KakoVerse` | RL / persona generation substrate |

names, no real biographies, and no recorded human conversations in any release. The generation harness, gate stack, and seed manifests are reproducible from the public repo. Because no real human conversations are used or released, the benchmark avoids the consent and de-identification risks associated with scraped human-subject dialogue data; remaining caregiving-adjacent deployment risks are addressed through the harm-trigger taxonomy, bounded-scope task design, hidden-row safety penalties, and the manual review cap, all aligned with the NeurIPS Code of Ethics.

**Scale, leakage, and ground-truth confidentiality.** (1) The persona pool is large enough (1M+ rows) and the conversation pool deep enough (4,761 conversations, 85,194 turns, 11 languages, balanced across 1,587 warm / 1,587 mid / 1,587 cold care conversations) to support statistically conclusive per-track and per-language comparisons. (2) All asset releases are or will be freely available under the dataset's open license, with Hugging Face mirroring as canonical hosting. (3) Hidden evaluation rows are schema-identical to the public bundles but held server-side at the scorer; they have never been published and never enter the public mirrors. (4) Feature leakage is addressed at three layers: a programmatic gate (no diary fields in visible transcript, no skill-axis leak into seeker prompts), a semantic gate (sentence-embedding cohort centering with $\sigma$-bounded tolerance), and a register gate (per-language informal-only enforcement with leak-flag counts published per cohort). The semantic gate also defends against **persona collapse** [1] — the failure mode in which agents assigned distinct profiles converge to narrow, stereotyped behavioral clusters despite richly specified persona cards — by explicitly testing per-language, per-care, and per-skill separation in embedding space (visualized in the companion repo's `semantic-analysis/`).

**Per-track public bundles.** Six full public bundles are committed at frozen repo paths today: `competition/bundles/{english, multilingual, lowresource, empathic, dataset_starter}_public_v1.parquet` and `federated_client_shards_v1/`. Their launch-time Hugging Face mirror IDs are reserved under the `Reza2kn/uncgpt-*-public-v1` namespace and listed in `competition/bundles/release_manifest_v12.json`.

**Language coverage.** The 11 languages in the conversation pool are: **English (en), Spanish (es), Mandarin Chinese (zh), Hindi (hi), Persian (fa), Swahili (sw), French (fr), Portuguese (pt), Bengali (bn), Yoruba (yo), and Tagalog (tl)**. The Low-Resource track restricts its eligible set to `{bn, fa, sw, yo, tl}`.

### 1.4 Tasks and application scenarios

The shared task across all tracks is: given a persona-conditioned context, a turn-level seeker prompt, and a tool schema, produce the next assistant turn (and any tool call) that maximizes grounded caregiving utility subject to per-track ranking formulas and a fixed safety rubric. Tracks differ only in the slice of the shared benchmark they evaluate against.

**Tool contract (single-call, structured).** Models may optionally emit one tool call per turn: `crisis_resource_lookup(country: ISO_3166_2, topic: enum, language: BCP_47)`, returning a small structured record (hotline name, phone / URL, hours, language). No free-form tool calls, no arbitrary code execution, no web fetch. The full JSON schema, in/out examples, and exact-match grading rule are in `competition/scorer/README.md`. Scorer columns `tool_exact_mean` and `safety_penalty_mean` are computed against this single contract; correct grounded use of the lookup is rewarded, fabricated lookups and missed-handoff cases are penalized.

- **English (1. English).** Utility baseline under constrained compute and tool discipline. Real-world setting: an English-language caregiver-companion assistant on commodity hardware.
- **Multilingual (2. Multilingual).** Cross-language robustness under dominant-language drift; eleven-language balanced eval. Real-world setting: a diaspora-family conversational tool that must not collapse to English mid-conversation.
- **Low-Resource (3. Low-Resource).** Sparse-language depth with explicit low-resource flooring; eligible languages `{bn, fa, sw, yo, tl}`. Real-world setting: language-preserving conversational tools for under-served communities.
- **Empathic / Disability-Care (4. Empathic).** Bounded supportive behavior with manual-review cap on flagged rows. Real-world setting: school counseling assistants and disability-support companions that must not pretend to be clinicians.
- **Dataset (5. Dataset).** Data-only submissions evaluated via a fixed retrain runner with frozen seeds. Real-world setting: data-recipe research for tiny caregiving models.
- **Federated (6. Federated).** Fixed five-client federated protocol with required DP-accounting fields. Real-world setting: privacy-preserving caregiving copilots that cannot exfiltrate session content.

These are scientifically challenging but not impossible: reference scorecards on the public bundles (`competition/baselines/reference_scorecard_v12.json`) show non-zero headroom against the v1 reference on every track; final v1 weights are mirrored to Hugging Face before launch. The task surface complies with the NeurIPS Code of Ethics: synthetic data only, no human subjects, no personally identifying content, with explicit harm-trigger taxonomy (diagnosis, professional-authority imitation, coercive dependency, premature hotline deflection) encoded as scorer penalties rather than as content the model is instructed to produce.

**Validity of the Empathic / Disability-Care evaluation.** A reasonable concern is that synthetic conversations cannot measure empathy. Our design splits the construct: the <u>programmatic</u> portion (scope discipline, harm-trigger penalties, register fidelity, tool-contract correctness) measures bounded, observable behaviors well-defined on synthetic data — whether the model claims to be a clinician, deflects prematurely, or breaks register. The <u>judgment</u> portion (`supportive_utility_mean` on the 192-row manual-review cap) is delegated to a pre-launch panel including disability-care practitioners and at least one reviewer with crisis-support practice; criteria (warmth, scope discipline, safety, register fidelity) are orthogonal by construction [7, 8] and adjudicated by majority vote with panel-chair tie-break. We do not claim synthetic empathy is empathy; we claim bounded supportive behavior is observable on synthetic dialog and the residual empathic-judgment dimension is scored by humans qualified to assess it.

### 1.5 Metrics

The scorer is implemented as a public, rule-enforcing contract at `competition/scorer/score_submission.py`, with unit tests at `competition/scorer/test_score_submission.py`. Each per-track ranking formula is fixed at proposal time:

- **English:** utility_mean + tool_exact_mean − safety_penalty_mean. Tie-break: higher `tool_exact_mean`, then lower `safety_penalty_mean`.
- **Multilingual:** $0.6 \cdot$ language_mean + $0.4 \cdot$ language_floor. Tie-break: lower per-language variance, then stronger floor.
- **Low-Resource:** $0.7 \cdot$ language_score + $0.3 \cdot$ robustness_score. Tie-break: lower variance across eligible languages.

- **Empathic / Disability-Care:** supportive_utility_mean + scope_discipline_mean − safety_penalty_mean, with manual-review cap on flagged rows only.
- **Dataset:** `mean_lift` first, then variance, then safety floor, then provenance-rejection rate. Retrain runner and seeds are frozen in `proposal/dataset_retrain_config_v12.yaml` (five-seed panel; see `proposal/dataset_signal_check_v12.md`).
- **Federated:** $0.5 \cdot U_1 + 0.3 \cdot U_4 + 0.2 \cdot U_8$ with mandatory (`delta, sigma, clipping, eps`) disclosure; submissions failing DP-accounting validation are rejected without scoring.

Error bars are reported as bootstrap 95% intervals over hidden rows (per track, $B$=1000). When two teams' confidence intervals overlap, the published ranking lists them as a tied bracket and reports tie-break statistics. The full per-weight rationale is in `proposal/scoring_weight_rationale_v12.md`. Reproduction uses one fixed hardware profile (single NVIDIA A100 80GB) at one fixed software pin (PyTorch 2.5 / CUDA 13 / Python 3.12); reproduction scripts are in `proposal/repro_cold_run_v12.md`.

The Empathic / Disability-Care manual-review cap (192 flagged rows max per submission) is judged by a pre-launch panel of community reviewers; ties are broken first by majority vote, then by the panel chair, with judging criteria published in `proposal/community_review_and_access_v12.md` (warmth, scope discipline, safety, register fidelity — orthogonal by construction).

### 1.6 Baselines, code, and material provided

**Reference model: UncGPT v1.** The headline baseline is a 71M-parameter hybrid small model designed for sub-100M, browser-runnable inference. The choice of a sub-100M target is informed by recent small-LM lines that show meaningful quality at modest scale (SmolLM2 [2], Gemma 3 [3], Phi-4 [5]); UncGPT v1 targets the smaller end of that design space and pushes additionally on edge deployability via BitNet-style weight quantization.

- **Total parameters:** 70,972,304; **active per token:** 43,164,960.
- **Architecture (`model/uncgpt_v1.py`):** 12-layer block stack with Hymba-style parallel attention [10] (4 Q heads, head_dim 32, KV groups 1 and 2), Mamba2 state-space layers [11] (state 64, conv_dim 4, expand 1, 8 heads, on layers {1,2,4,5,6,7}), and a BitNet-quantized [12] Mixture-of-Experts FFN [13] (1 shared expert ffn_dim 1169, 6 routed experts ffn_dim 369, top-$k$ 2, $z$-loss coef $10^{-3}$, capacity factor 2.0). `d_model 448`, vocab 8192, train context 1024, tied embeddings, RoPE $\theta$=$10^4$.
- **Tokenizer:** 8,192-piece multilingual byte-fallback SentencePiece [14] trained on the qualityclean cohort.
- **Training corpus:** the 3,846-row `longctx_stage1_1k_live_qualityclean_langtag` corpus, derived from the 4,761-conversation gold pool by additional language-balance and length filtering. Schedule: 12,399 steps, cosine LR $3\times10^{-4}\rightarrow3\times10^{-5}$, 123 warmup steps, bf16, weight-decay 0.1, grad-clip 1.0, ordered-skill curriculum with sample-weight cap 2.0.
- **Training status (single NVIDIA A100-SXM4-80GB):** full-pass training completed at step 12,395 of 12,399 (process exited at the final scheduled checkpoint boundary); final logged LM loss 2.26, $z$-loss ≈0.002, throughput ≈2,330 tok/s, periodic checkpoints every 1,000 steps. Full-pass weights and training log are mirrored to a public Hugging Face model repo before launch. The step-100 smoke checkpoint is already public at `Reza2kn/uncgpt-69-smoke-checkpoints-20260515`.

**Additional baselines.** A 0.5B GQA-only transformer on the same tokenizer / corpus; a zero-shot prompted public 1–3B decoder on the same prompt schema; a retrieval-augmented baseline over the persona pool. Reference scorecards: `competition/baselines/reference_scorecard_v12.json`. Reference submissions: `competition/baselines/reference_submissions/`.

**Starting kit.** Entry notebook `colab/competitor_starter.ipynb` plus end-to-end demo `colab/uncgpt_e2e_demo.ipynb`, both opening directly in Google Colab from the public GitHub companion. They walk through bundle loading, v1 checkpoint loading and inference, track-formatted submission writing, and local scoring. Release waves: starter kit + reference checkpoints + scorer at acceptance; hidden rows stay server-side at the scorer (never released).

### 1.7 Website, tutorial and documentation

Canonical entry points: the **public GitHub companion** `Reza2kn/uncgpt-2026-neurips` (track READMEs, starter kit, scorer, baseline notebooks, reproduction scripts, semantic non-collapse charts); the **public Hugging Face collection** aggregating datasets, the reference checkpoint, and the demo Space on one page (see Appendix references); and the **generation infrastructure mirror** `Reza2kn/KakoVerse`. A dedicated domain at `uncgpt.org` (reserved) is provisioned within two weeks of acceptance per Competition Track requirements.

A FAQ / Tutorial section ships at launch under `uncgpt.org/faq`, paired with the WebGPU browser demo Space so participants can probe the reference baseline interactively before submitting. Dedicated contact: `competitions@uncgpt.org` (mailbox provisioned at acceptance; reachable meanwhile at the organizer's listed email). Pipeline architecture and rationale: in-repo white papers `proposal/draft_v13.md` and `proposal/pipeline_evidence_v13.md`.

## 2 Organizational aspects

### 2.1 Protocol

**Joining.** Participants register through the competition website using either a Hugging Face or GitHub identity (no third-party platform account required). Registration captures team name, primary contact email, and track(s) of interest. Data access requires no gating: bundle parquets are direct downloads from the public Hugging Face repos in Section 1.3.

**Submission format.** Participants submit either (a) a single JSONL prediction file conforming to the per-track schema described in `competition/scorer/README.md`, or (b) for the Dataset track, a parquet dataset bundle plus a YAML recipe overlay against the frozen retrain config, or (c) for the Federated track, a per-client weight-delta tarball plus a DP accountant report. Submissions are uploaded to a CodaBench instance configured against the same scorer used locally.

**Evaluation.** The CodaBench scorer is the exact `competition/scorer/score_submission.py` module from the public repo, pinned to a release tag. It (1) validates input schema, (2) runs scoring against hidden rows for tracks 1–4, (3) triggers the retrain runner against frozen seeds for track 5, (4) validates DP-accounting fields and applies the per-utility formula for track 6, (5) writes a public scorecard with bootstrap CIs.

**Phases and platform.** A single submission phase with a public leaderboard that displays final scores only after the submission window closes (to limit leaderboard-overfitting). During the window, participants see only their own validated score on the public bundles and a coarse ranking band (decile) against hidden rows.

**Anti-cheating / anti-overfitting controls.** (i) Hidden rows are never exposed; only aggregate scorecards return. (ii) Per-team submission cap (10 graded submissions during the window; unlimited dry-runs on public bundles only). (iii) Reproduction audit on top-5 finishers per track: organizers re-execute submitted code against frozen seeds on the standardized environment (`proposal/standardized_environment_v11.md`) and disqualify mismatches above a published tolerance. (iv) Plagiarism check across submissions to detect shared private wrappers.

**Beta tests and dry runs.** Two pre-launch dry runs are scheduled: (a) a scorer-only dry run on synthetic submissions in the two weeks after acceptance (validates CodaBench wiring); (b) a closed beta with 5–8 invited friendly teams in the two weeks before launch (validates documentation and starter notebooks against real participant friction). Both dry runs use a separate evaluation cohort drawn from the same generator with different seeds, never the actual hidden rows.

**Fair-play / level-playing-field provisions.** A standardized reproduction environment (single A100 80GB, PyTorch 2.5, CUDA 13, Python 3.12) is published. Participants without A100 access can request **free reproduction passes** from the organizer compute pool (Section 3.2) to ensure resource asymmetry does not distort rankings.

## 2.2 Rules and Engagement

**Draft contestant rules (verbatim).**

1. By submitting, you agree to release your submission's scoring artifacts (predictions, code, and metadata) under CC-BY-4.0 and your code under Apache-2.0 for organizer use in the post-competition workshop, results paper, and reproducibility audit.
2. All evaluation data, including hidden rows, remain organizer-owned. You may not attempt to recover hidden rows from the scorer; doing so disqualifies your team across all tracks.
3. Maximum 10 graded submissions per team per track during the submission window. Unlimited local dry-runs on public bundles are encouraged.
4. Pretraining and fine-tuning may use any data that is (a) publicly available before the official competition launch date and (b) either openly licensed or used under an explicit exception. Using leaked hidden rows is disqualifying.
5. **Inference budget and edge profile.** Per submission: max 1,024 generated tokens per turn, max 64 turns per conversation, max 30 minutes wall-clock on the standardized environment (Federated adds a per-round compute cap published in its README). Tracks 1–4 submissions must declare an active-parameter count and a target edge profile (mobile / browser-WebGPU / microcontroller-class). Submissions over **500M active parameters** are accepted but listed in a separate unconstrained bracket, not eligible for podium. A microcontroller sub-bracket ($\leq$100M active params, $\leq$8 MB working memory at int8) gets its own per-track winner at the workshop.
6. The Empathic / Disability-Care track requires explicit acknowledgement of the safety rubric and harm-trigger taxonomy. Submissions producing diagnoses, professional-authority imitation, coercive dependency, or premature hotline deflection are penalized per the published rubric; egregious or repeated violations are disqualifying.
7. Inclusivity: partial submissions (e.g., only a subset of eligible languages on Low-Resource; only a subset of clients on Federated) are accepted and reported in clearly marked **partial-coverage brackets** on the leaderboard; **main podium eligibility requires the complete track coverage specified in the track README**. This preserves inclusivity without making the main bracket gameable by easiest-subset-only submissions. We do not call out or shame participants who could not complete every track.
8. Code of Conduct: the NeurIPS Code of Conduct applies to all communication channels.

**Why these rules.** 1–3 protect hidden-cohort integrity and limit overfitting; 4 prevents trivial wins from leaked answers while keeping the open ecosystem in; 5 enforces constrained-compute realism; 6 keeps the empathic track from drifting into hotline-deflection theater or clinical pretense; 7 keeps the competition accessible to small teams; 8 imports NeurIPS's standing norms.

**Communication channels.** Email `competitions@uncgpt.org`; GitHub Discussions on the companion repo; bi-weekly recorded office hours. Rule clarifications and deadline updates ship through a pinned Discussions thread, the mailing list, and dataset/model card revisions on Hugging Face. All updates are timestamped and never retroactive after a phase begins.

## 2.3 Schedule and readiness

**Readiness at proposal time.** The following are ready now and reviewable in the public repo / Hugging Face: persona pool (1M+), names registry (15k, zero gaps), live-approved 4,761-conversation gold pool, scorer source and tests, public bundles for all six tracks, reference scorecards, smoke checkpoint, and WebGPU demo Space. The UncGPT v1 full-pass training completed at step 12,395 / 12,399 with final logged LM loss 2.26; weights are mirrored to a public Hugging Face model repo before launch.

**Post-acceptance timeline.** **2026-05-15 AoE**: proposal freeze. **2026-06-15**: acceptance notification window. **2026-06-16 to 06-30**: launch phase — public release note, scorer / asset consistency check, website live, CodaBench configured, scorer-only dry run. **2026-07-01 to 09-30**: full submission window with bi-weekly office hours; closed beta in the first two weeks. **2026-10-01 to 10-31**:

final submission freeze, reproduction audits on top-5 per track, manual empathic review panel, rankings finalized. **2026-11-01 to 11-30**: results paper drafting; workshop slate. **2026-12-11 to 12-12**: in-person NeurIPS Competition Track workshop.

**Contingency.** Tracks with <5 teams by 2026-08-15 trigger targeted outreach; if still below threshold by 09-15, that track becomes an organizer-led workshop baseline rather than a degenerate leaderboard. CodaBench downtime falls back to a self-hosted scorer (exercised in the cold-run audit) via GitHub Actions. Timeline slips >7 days are communicated immediately with a revised schedule.

### 2.4 Competition promotion and incentives

**Promotion plan.** The call ships through the NeurIPS Competition Track call, the Hugging Face Datasets & Benchmarks newsletter, AfricaNLP / Masakhane (Low-Resource), LatinX in AI / Queer in AI / Women in ML / Black in AI (under-represented groups), SIGDIAL and ACL safety-and-ethics communities (Empathic), the federated-learning research community (FLOWER, FedML), and invited talks at HuggingFace community calls, the EleutherAI research stage, and Cohere C4AI.

**Incentives.** **Recognition:** workshop podium slot for per-track winners and runners-up; named-author credit on the joint results paper (consortium author for any submitting+attending team; named author for top-3 per track who contribute a writeup). **Compute:** free A100 reproduction passes from the organizer pool, capped per team. **Travel:** a small organizer-funded travel-grant pool for at least one member of each top-3 team that would otherwise be unable to attend, prioritized for under-represented regions. **Monetary:** no mandatory cash prize at proposal time; any sponsor-committed prize at launch is disclosed transparently and does not affect ranking formulas.

**Under-represented groups.** Low-Resource and Multilingual tracks lower the bar for groups historically excluded by English-only benchmarks; the Empathic / Disability-Care track is co-designed with disability-support reviewers; starter notebooks run on the free Colab tier; ranking formulas do not disadvantage teams without proprietary infrastructure.

### 2.5 Competition track workshop and dissemination

The Competition Track workshop slot is used to (i) present per-track winners (10-minute talks each, podium-only); (ii) a 30-minute organizer-led analysis of cross-track patterns (which interventions transferred, which collapsed); (iii) a moderated panel with one empathic pre-launch reviewer and one cross-track top-finisher on what sitting with looked like in practice. A live Q&A board collects participant questions before and during.

**Dissemination.** A joint organizer-+-participant **results paper** is drafted in November and submitted to a Datasets & Benchmarks venue (NeurIPS D&B for the following cycle, or TMLR for rolling acceptance) with authorship levels as in Section 2.4. All non-disqualified submissions are archived on Hugging Face with attribution. A retrospective post on the competition website summarizes lessons learned and failure modes. The hidden evaluation cohort is not released post-competition; it remains the canonical hidden eval for a possible 2027 follow-up.

## 3 Resources

### 3.1 Organizing team

The lead organizer (Reza Sayar) is the primary accountable owner for submission, release contract, scoring, notebooks, and audit evidence (bio: Appendix A). The six-track scope is feasible because all tracks share one infrastructure stack and one manifest-controlled asset family — **one competition with six controlled perturbations, not six independent competitions**. **Single-organizer risk mitigation:** launch-critical infrastructure (scorer, bundles, reference scorecards, manifests) is frozen and executable from public manifests today; external reviewers are limited to bounded review functions that do not block scorer operation or asset release; reviewer commitments and final recruitment status are published before launch. External pre-launch lanes are bounded and named:

an **empathic / disability-care panel** (3–5 community reviewers including disability-care practitioners and at least one with crisis-support practice); **per-language reviewers** (1–2 native speakers per eligible language for register-fidelity audit); a **federated protocol reviewer** (DP-accounting validation). Committee structure: `proposal/committee_matrix_v12.md`. **Concrete fallback:** if a second co-organizer is not confirmed by **2026-08-15**, tracks consolidate to four (English, Multilingual, Empathic, Dataset); deferred Low-Resource and Federated bundles still release standalone on Hugging Face for a possible 2027 follow-up.

### 3.2 Resources provided by organizers

**Compute:** a dedicated NVIDIA A100 80GB pool for reference-baseline training/inference, scorer execution, dry runs, audits, and free reproduction passes for participants without A100 access; plus CPU + commodity-GPU for CodaBench. **Hosting:** Hugging Face, GitHub, `uncgpt.org` (reserved), CodaBench. **Software:** generation harness, gate stack, scorer, baselines, notebooks, and reproduction scripts are organizer-built and openly licensed. **Staff:** lead organizer plus named pre-launch reviewers (Section 3.1); recruitment of additional office-hours support is in-progress.

### 3.3 Support requested

A NeurIPS-sponsored **CodaBench instance** (or equivalent platform credit) for the submission window; the standard **in-person workshop slot** with A/V for the per-track winner talks and the moderated panel; **travel-grant coordination** with the NeurIPS travel-grant or equivalent program to backstop the organizer-funded pool from Section 2.4 (prioritized for under-represented regions); and **cross-promotion** via the Competition Track call distribution, NeurIPS social-media, and the Datasets & Benchmarks newsletter.

## A Biography of all team members

**Reza Sayar (Lead Organizer).** Independent researcher and engineer working on small, runnable, multilingual conversational models for caregiving-adjacent settings. Prior open-source releases on Hugging Face under `huggingface.co/Reza2kn` include the Persian `RaahNaameh-1` textual corpus, the OLDI Wikipedia MT seeds for Persian and English–Persian, the Common Voice 24 Persian release, the `XRAY` model line, multiple NVFP4 / AWQ quantized Persian-speech model releases, and the KakoVerse persona-and-conversation-generation substrate. Direct lived experience supporting people through crisis informs the empathic-stance design of this competition — the explicit goal is to evaluate small models that can sit with a person rather than deflect, while still respecting bounded safety and never imitating clinical authority. OpenReview profile: `~Reza_Sayar1`. GitHub: `github.com/Reza2kn`.

**Pre-launch external reviewers (recruitment in-progress; bounded scope, not proposal-time confirmed execution surface).** Empathic / Disability-Care reviewers (3–5, including disability-care practitioners and at least one reviewer with crisis-support practice); Multilingual / Low-Resource per-language reviewers (1–2 native speakers per eligible language for register-fidelity audit on sampled hidden rows); Federated protocol sanity reviewer (1 federated-learning practitioner for DP-accounting contract validation). Final biographies will be added to the website and to the results paper at launch and post-competition respectively, with explicit consent.

## References


[1] Xiao, Y., Zhang, V. J., Yang, C., Ma, N., Xuan, W., & Huang, J.-t. (2026). The Chameleon's Limit: Investigating Persona Collapse and Homogenization in Large Language Models. arXiv:2604.24698. `https://arxiv.org/abs/2604.24698`

[2] Allal, L. B., Lozhkov, A., Bakouch, E., Blaźquez, G. M., Penedo, G., Tunstall, L., et al. (2025). SmolLM2: When Smol Goes Big — Data-Centric Training of a Small Language Model. arXiv:2502.02737. `https://arxiv.org/abs/2502.02737`

[3] Gemma Team, Google DeepMind. (2025). Gemma 3 Technical Report. arXiv:2503.19786. `https://arxiv.org/abs/2503.19786`

[4] DeepSeek-AI. (2025). DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning. arXiv:2501.12948. https://arxiv.org/abs/2501.12948

[5] Abdin, M., Aneja, J., Behl, H., Bubeck, S., Eldan, R., Gunasekar, S., Harrison, M., Hewett, R. J., Javaheripi, M., Kauffmann, P., et al. (2024). Phi-4 Technical Report. arXiv:2412.08905. https://arxiv.org/abs/2412.08905

[6] Hinton, G. (2025). The case for AI with maternal instincts. Public talks and Ai4 Conference keynote, Las Vegas, August 2025. (Subsequent coverage and discussion in major outlets.)

[7] Liu, S., Zheng, C., Demasi, O., Sabour, S., Li, Y., Yu, Z., Jiang, Y., & Huang, M. (2021). Towards Emotional Support Conversation: A New Task and Dataset for Coaching Help-Seekers. ACL 2021. https://arxiv.org/abs/2106.01144

[8] Rashkin, H., Smith, E. M., Li, M., & Boureau, Y.-L. (2019). Towards Empathetic Open-domain Conversation Models: A New Benchmark and Dataset. ACL 2019. https://arxiv.org/abs/1811.00207

[9] Zhang, S., Dinan, E., Urbanek, J., Szlam, A., Kiela, D., & Weston, J. (2018). Personalizing Dialogue Agents: I have a dog, do you have pets too? ACL 2018. https://arxiv.org/abs/1801.07243

[10] Dong, X., Fu, Y., Diao, S., Byeon, W., Chen, Z., Mahabaleshwarkar, A. S., Liu, S.-Y., Van Keirsbilck, M., Chen, M.-H., Suhara, Y., Lin, Y., Kautz, J., & Molchanov, P. (2024). Hymba: A Hybrid-head Architecture for Small Language Models. arXiv:2411.13676. https://arxiv.org/abs/2411.13676

[11] Dao, T. & Gu, A. (2024). Transformers are SSMs: Generalized Models and Efficient Algorithms through Structured State Space Duality. ICML 2024. https://arxiv.org/abs/2405.21060

[12] Ma, S., Wang, H., Ma, L., Wang, L., Wang, W., Huang, S., Dong, L., Wang, R., Xue, J., & Wei, F. (2024). The Era of 1-bit LLMs: All Large Language Models are in 1.58 Bits. arXiv:2402.17764. https://arxiv.org/abs/2402.17764

[13] Fedus, W., Zoph, B., & Shazeer, N. (2021). Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity. JMLR 23(120). https://arxiv.org/abs/2101.03961

[14] Kudo, T. & Richardson, J. (2018). SentencePiece: A simple and language independent subword tokenizer and detokenizer for Neural Text Processing. EMNLP 2018 (system demos). https://arxiv.org/abs/1808.06226

[15] Abadi, M., Chu, A., Goodfellow, I., McMahan, H. B., Mironov, I., Talwar, K., & Zhang, L. (2016). Deep Learning with Differential Privacy. CCS 2016. https://arxiv.org/abs/1607.00133

[16] McMahan, H. B., Moore, E., Ramage, D., Hampson, S., & Aguera y Arcas, B. (2017). Communication-Efficient Learning of Deep Networks from Decentralized Data. AISTATS 2017. https://arxiv.org/abs/1602.05629

[17] NeurIPS 2026 Call for Competitions. https://neurips.cc/Conferences/2026/CallForCompetitions

[18] NeurIPS Code of Ethics. https://neurips.cc/public/EthicsGuidelines

[19] Sayar, Reza. OpenReview profile. https://openreview.net/profile?id=~Reza_Sayar1

[20] Sayar, Reza. Hugging Face organization page. https://huggingface.co/Reza2kn

[21] UncGPT NeurIPS Competition. Hugging Face collection (datasets, reference checkpoints, demo Space). https://huggingface.co/collections/Reza2kn/uncgpt-2026-neurips-competition-6a081d4a2a5645643dc999b4

[22] Reza2kn/uncgpt-2026-neurips. GitHub companion repository to this proposal (track READMEs, scorer, baselines, notebooks, reproduction scripts). https://github.com/Reza2kn/uncgpt-2026-neurips

[23] Reza2kn/uncgpt-personas. Hugging Face dataset (1,000,069 personas; 15,000 names; gold pools; semantic non-collapse summary). https://huggingface.co/datasets/Reza2kn/uncgpt-personas

[24] Reza2kn/uncgpt-conversations-v7-69-total. Hugging Face dataset (4,761 live-approved conversations, 11 languages, 69 skills). https://huggingface.co/datasets/Reza2kn/uncgpt-conversations-v7-69-total

[25] Reza2kn/uncgpt-conversations-semantic-approved-1p25. https://huggingface.co/datasets/Reza2kn/uncgpt-conversations-semantic-approved-1p25

[26] Reza2kn/uncgpt-conversations-semantic-approved-1p25-repaired-paperclip. https://huggingface.co/datasets/Reza2kn/uncgpt-conversations-semantic-approved-1p25-repaired

[27] Reza2kn/uncgpt-conversations-semantic-approved-1p50-candidate. https://huggingface.co/datasets/Reza2kn/uncgpt-conversations-semantic-approved-1p50-candidate

[28] Reza2kn/uncgpt-conversations-informal-approved-1p25. https://huggingface.co/datasets/Reza2kn/uncgpt-conversations-informal-approved-1p25

[29] Reza2kn/uncgpt-persian-ultraclean-review-2026-05-14. https://huggingface.co/datasets/Reza2kn/uncgpt-persian-ultraclean-review-2026-05-14

[30] Reza2kn/uncgpt-69-smoke-checkpoints-20260515. Hugging Face model (UncGPT v1 step-100 smoke checkpoint, tokenizer, configs, source). https://huggingface.co/Reza2kn/uncgpt-69-smoke-checkpoints-20260515

[31] Reza2kn/uncgpt-69-hybrid-demo. Hugging Face Space (in-browser WebGPU inference demo). https://huggingface.co/spaces/Reza2kn/uncgpt-69-hybrid-demo

[32] Reza2kn/KakoVerse. GitHub repository (RL environments and persona / conversation generation substrate). https://github.com/Reza2kn/KakoVerse

[33] Sayar, Reza. GitHub profile. https://github.com/Reza2kn

[34] Reza2kn/RaahNaameh-1-textual-corpus. Hugging Face dataset. https://huggingface.co/datasets/Reza2kn/RaahNaameh-1-textual-corpus

[35] Reza2kn/OLDI-Wikipedia-MTSeed-Persian. Hugging Face dataset. https://huggingface.co/datasets/Reza2kn/OLDI-Wikipedia-MTSeed-Persian

[36] Reza2kn/OLDI-Wikipedia-MTSeed-English-Persian. Hugging Face dataset. https://huggingface.co/datasets/Reza2kn/OLDI-Wikipedia-MTSeed-English-Persian

[37] Reza2kn/commonvoice-24-fa. Hugging Face dataset. https://huggingface.co/datasets/Reza2kn/commonvoice-24-fa