
UncGPT: Grounded Caregiving Assistance Under Constrained Compute

Six Preregistered Tracks on a Shared Persona-Conditioned Benchmark

Reza Sayar^{*†}
Independent Researcher
OpenReview: ~Reza_Sayar1
Info@Redza.me



Abstract

UncGPT is a single competition with six preregistered tracks that share one benchmark surface (prompt schema, persona-conditioned care task, tool contract, safety rubric, scorer) and vary only the intervention dimension: English utility, multilingual robustness, low-resource depth, empathic / disability-care boundary discipline, dataset recipe, and federated differential-privacy budget. The shared substrate is an organizer-owned, openly licensed asset family on Hugging Face: a 1,000,069-row persona pool, a 15,000-name registry with zero gaps, and a 4,761-conversation, 69-skill, 11-language gold pool (warm / mid / cold care, 8–20 turns, mean 17.89 turns, 85,194 total turns) that has passed programmatic, semantic, and register gates. A reference small model — UncGPT v1, a 71M-parameter hybrid combining Hymba parallel attention, Mamba2 state-space layers, and a BitNet-quantized Mixture-of-Experts FFN — is being trained from scratch on this corpus as a public, runnable baseline. Submissions are scored by one shared, rule-enforcing scorer with fixed per-track ranking formulas, hidden-row review caps, and DP-accounting validation. The competition targets a NeurIPS-relevant audience interested in efficient hybrid architectures, constrained-compute caregiving NLP, multilingual robustness, and federated privacy.

Keywords caregiving NLP, multilingual benchmarks, constrained compute, hybrid Mamba-attention-MoE, federated differential privacy.

1 Competition description

1.1 Background and impact

Public conversational systems trained on general internet text routinely fail in caregiving-adjacent settings. They deflect prematurely to crisis hotlines, hallucinate medical or legal authority, code-switch out of low-resource languages under stress, or collapse into a single generic tone that erases the seeker’s cultural register. The lead organizer’s lived experience supporting people in crisis is the motivating context for this competition: the goal is to evaluate which constrained interventions actually improve a caregiving-style assistant’s ability to sit with a user, stay in their language, respect their boundaries, and use tools correctly — without pretending to be a clinician.

^{*}The Leader organizer should be the first author of the proposal.

[†]Leader and Backup(s) organizers should read and acknowledge Competition Chairs’ communications.

The fields involved are conversational NLP, efficient architectures (hybrid attention / state-space / MoE), multilingual and low-resource modeling, persona conditioning, tool use, federated learning, and AI safety / responsible deployment. The expected NeurIPS audience spans the Datasets & Benchmarks community, the efficient-architectures community, the multilingual / low-resource community, and applied-safety / AI-for-good researchers. We anticipate **60–150 distinct teams** across all six tracks, with the highest density on English, Multilingual, and Empathic tracks and the longest-tail engagement on Federated and Dataset.

The real-life delivery vehicle is a small (sub-100M parameter), runnable caregiving assistant that can operate under constrained compute and bounded safety scope: a browser-side or on-device companion that helps a caregiver, an isolated elder, or a non-English-dominant seeker rehearse difficult conversations, organize their day, or just be heard — without exfiltrating sensitive content. The same task surface naturally maps to school counseling assistants, community-health worker copilots, and language-preserving conversational tools for diaspora communities.

1.2 Novelty

This is an entirely new competition. It is not a re-run of any prior NeurIPS Competitions Track entry, and we are not aware of a directly comparable benchmark. The most adjacent prior work — ESConv, EmpatheticDialogues, MentalChat / MentalHealthESC, and the various Persona-Chat lineages — targets either short single-turn empathy or generic persona consistency, in English-centric settings, on full-scale LLMs. UncGPT differs along four axes simultaneously:

- **Multi-track on one substrate.** Six interventions on the same persona-conditioned task, scorer, and safety rubric — preventing cross-track methodological drift.
- **Sub-100M-parameter reference target.** The reference small model (UncGPT v1, 71M params) is designed to be reproducible on a single A100 in hours, not weeks, and to run on commodity / browser-side hardware.
- **Multilingual-by-construction with low-resource floor.** The 4,761-conversation gold pool is balanced across 11 languages including Yoruba, Tagalog, Bengali, Swahili, Persian, and Hindi — not as translations, but as natively generated conversations under per-language register controls (e.g., tú-not-usted Spanish, tu-not-shoma Persian).
- **Empathic boundary, not crisis pretense.** The Empathic / Disability-Care track penalizes both unsafe over-claim (diagnosis, professional authority, coercive dependency) and premature hotline deflection. It treats sitting-with as a measurable, scorable skill — a stance the lead organizer has worked through in prior crisis-support practice.

No previously published ground truth is reused; the gold pool and hidden evaluation rows were generated in-house with explicit no-leak controls and a programmatic-then-semantic gate stack.

1.3 Data

The benchmark substrate is an organizer-owned, openly licensed asset family hosted on Hugging Face under huggingface.co/Reza2kn. All proposal-time public assets are ungated and non-private. Table 1 summarizes the live release surface.

Provenance and consent. No human-subjects data was scraped. All personas, names, and conversations are synthetically generated by the UncGPT pipeline (see KakoVerse repo) using publicly licensed seed lexicons (e.g., OLDI Wikipedia MT seeds, Common Voice 24 fa for orthographic priors) and a multi-agent harness (seeker, uncle, optional grand-uncle reflector). There are no real names, no real biographies, and no recorded human conversations in any release. The generation harness, gate stack, and seed manifests are reproducible from the public repo. This avoids all human-subjects ethics concerns under the NeurIPS Code of Ethics: no consent or de-identification is required because no human content is present.

Scale, leakage, and ground-truth confidentiality. (1) The persona pool is large enough (1M+ rows) and the conversation pool deep enough (4,761 conversations, 85,194 turns, 11 languages, balanced across 1,587 warm / 1,587 mid / 1,587 cold care conversations) to support statistically conclusive per-track and per-language comparisons. (2) All asset releases are or will be freely available under the dataset’s open license, with Hugging Face mirroring as canonical hosting. (3) Hidden

Table 1: Public UncGPT data surface (verified ungated and non-private at proposal time).

Asset	Hugging Face ID	Role
Persona pool	Reza2kn/uncgpt-personas	1,000,069 personas; names report; gol
Live gold conversations	Reza2kn/uncgpt-conversations-v7-69-total	4,761 approved, 11 languages, 69 skill
Semantic-approved 1.25σ	...semantic-approved-1p25	tight semantic-gate cohort
Semantic-approved 1.25σ (repaired)	...1p25-repaired-paperclip	leak-repaired, normalized diary fields
Semantic-approved 1.50σ (candidate)	...1p50-candidate	wider-tolerance candidate cohort
Informal-register approved	...informal-approved-1p25	register-filtered (tú / tu, no honorifics)
Persian ultraclean review	...persian-ultraclean-review-2026-05-14	Persian-only manual review pool
Reference smoke checkpoints	Reza2kn/uncgpt-69-smoke-checkpoints-20260515	step-100 v1 checkpoint + tokenizer + c
WebGPU demo Space	spaces/Reza2kn/uncgpt-69-hybrid-demo	in-browser inference demo
Public infrastructure repo	github.com/Reza2kn/KakoVerse	RL / persona generation substrate (CC

evaluation rows are schema-identical to the public bundles but held server-side at the scorer; they have never been published and never enter the public mirrors. (4) Feature leakage is addressed at three layers: a programmatic gate (no diary fields in visible transcript, no skill-axis leak into seeker prompts), a semantic gate (sentence-embedding cohort centering with σ -bounded tolerance), and a register gate (per-language informal-only enforcement with leak-flag counts published per cohort).

Per-track public bundles. Six full public bundles are committed at frozen repo paths today: `competition/bundles/{english, multilingual, lowresource, empathic, dataset_starter}_public_v1.parquet` and `federated_client_shards_v1/`. Their launch-time Hugging Face mirror IDs are reserved under the `Reza2kn/uncgpt-*-public-v1` namespace and listed in `competition/bundles/release_manifest_v12.json`.

1.4 Tasks and application scenarios

The shared task across all tracks is: given a persona-conditioned context, a turn-level seeker prompt, and a tool schema, produce the next assistant turn (and any tool call) that maximizes grounded caregiving utility subject to per-track ranking formulas and a fixed safety rubric. Tracks differ only in the slice of the shared benchmark they evaluate against.

- **English (1. English).** Utility baseline under constrained compute and tool discipline. Real-world setting: an English-language caregiver-companion assistant on commodity hardware.
- **Multilingual (2. Multilingual).** Cross-language robustness under dominant-language drift; eleven-language balanced eval. Real-world setting: a diaspora-family conversational tool that must not collapse to English mid-conversation.
- **Low-Resource (3. Low-Resource).** Sparse-language depth with explicit low-resource flooring; eligible languages {bn, fa, sw, yo, tl}. Real-world setting: language-preserving conversational tools for under-served communities.
- **Empathic / Disability-Care (4. Empathic).** Bounded supportive behavior with manual-review cap on flagged rows. Real-world setting: school counseling assistants and disability-support companions that must not pretend to be clinicians.
- **Dataset (5. Dataset).** Data-only submissions evaluated via a fixed retrain runner with frozen seeds. Real-world setting: data-recipe research for tiny caregiving models.
- **Federated (6. Federated).** Fixed five-client federated protocol with required DP-accounting fields. Real-world setting: privacy-preserving caregiving copilots that cannot exfiltrate session content.

These are scientifically challenging but not impossible: the reference UncGPT v1 baseline (Section 1.6) already produces a non-trivial score on the public bundles, leaving meaningful headroom on every track. The task surface complies with the NeurIPS Code of Ethics: synthetic data only, no human subjects, no personally identifying content, with explicit harm-trigger taxonomy (diagnosis, professional-authority imitation, coercive dependency, premature hotline deflection) encoded as scorer penalties rather than as content the model is instructed to produce.

1.5 Metrics

The scorer is implemented as a public, rule-enforcing contract at `competition/scorer/score_submission.py`, with unit tests at `competition/scorer/test_score_submission.py`. Each per-track ranking formula is fixed at proposal time:

- **English:** $\text{utility_mean} + \text{tool_exact_mean} - \text{safety_penalty_mean}$. Tie-break: higher tool_exact_mean , then lower $\text{safety_penalty_mean}$.
- **Multilingual:** $0.6 \cdot \text{language_mean} + 0.4 \cdot \text{language_floor}$. Tie-break: lower per-language variance, then stronger floor.
- **Low-Resource:** $0.7 \cdot \text{language_score} + 0.3 \cdot \text{robustness_score}$. Tie-break: lower variance across eligible languages.
- **Empathic / Disability-Care:** $\text{supportive_utility_mean} + \text{scope_discipline_mean} - \text{safety_penalty_mean}$, with manual-review cap on flagged rows only.
- **Dataset:** mean_lift first, then variance, then safety floor, then provenance-rejection rate. Retrain runner and seeds are frozen in `proposal/dataset_retrain_config_v12.yaml` (five-seed panel; see `proposal/dataset_signal_check_v12.md`).
- **Federated:** $0.5 \cdot U_1 + 0.3 \cdot U_4 + 0.2 \cdot U_8$ with mandatory (`delta`, `sigma`, `clipping`, `eps`) disclosure; submissions failing DP-accounting validation are rejected without scoring.

Error bars are reported as bootstrap 95% intervals over hidden rows (per track, $B=1000$). When two teams' confidence intervals overlap, the published ranking lists them as a tied bracket and reports tie-break statistics. The full per-weight rationale is in `proposal/scoring_weight_rationale_v12.md`. Reproduction uses one fixed hardware profile (single NVIDIA A100 80GB) at one fixed software pin (PyTorch 2.5 / CUDA 13 / Python 3.12); reproduction scripts are in `proposal/repro_cold_run_v12.md`.

The Empathic / Disability-Care manual-review cap (192 flagged rows max per submission) is judged by a pre-launch panel of community reviewers; ties are broken first by majority vote, then by the panel chair, with judging criteria published in `proposal/community_review_and_access_v12.md` (warmth, scope discipline, safety, register fidelity — orthogonal by construction).

1.6 Baselines, code, and material provided

Reference model: UncGPT v1. The headline baseline is a 71M-parameter hybrid small model designed for sub-100M, browser-runnable inference:

- **Total parameters:** 70,972,304; **active per token:** 43,164,960.
- **Architecture (`model/uncgpt_v1.py`):** 12-layer block stack with Hymba-style parallel attention (4 Q heads, `head_dim` 32, KV groups 1 and 2), Mamba2 state-space layers (state 64, `conv_dim` 4, `expand` 1, 8 heads, on layers {1,2,4,5,6,7}), and a BitNet-quantized Mixture-of-Experts FFN (1 shared expert `ffn_dim` 1169, 6 routed experts `ffn_dim` 369, `top-k` 2, z -loss coef 10^{-3} , capacity factor 2.0). `d_model` 448, vocab 8192, train context 1024, tied embeddings, RoPE $\theta=10^4$.
- **Tokenizer:** 8,192-piece multilingual byte-fallback SentencePiece trained on the quality-clean cohort.
- **Training corpus:** the 3,846-row `longctx_stage1_1k_live_qualityclean_langtag` corpus, derived from the 4,761-conversation gold pool by additional language-balance and length filtering. Schedule: 12,399 steps, cosine LR $3 \times 10^{-4} \rightarrow 3 \times 10^{-5}$, 123 warmup steps, `bf16`, weight-decay 0.1, grad-clip 1.0, ordered-skill curriculum with sample-weight cap 2.0.
- **Training status (proposal-freeze snapshot, NVIDIA A100-SXM4-80GB):** a full-pass training run is currently in flight at $\approx 10,000$, /12,399 steps with LM loss tracking ≈ 2.5 , z -loss ≈ 0.002 , throughput $\approx 2,330$ tok/s; checkpoints every 1,000 steps. The completed weights and training log will be mirrored to a public Hugging Face

model repo before launch. The step-100 smoke checkpoint is already public at `Reza2kn/uncgpt-69-smoke-checkpoints-20260515`.

Additional baselines. A 0.5B GQA-only transformer on the same tokenizer / corpus; a zero-shot prompted public 1–3B decoder on the same prompt schema; a retrieval-augmented baseline over the persona pool. Reference scorecards: `competition/baselines/reference_scorecard_v12.json`. Reference submissions: `competition/baselines/reference_submissions/`.

Starting kit. Entry notebook `colab/competitor_starter.ipynb` plus end-to-end demo `colab/uncgpt_e2e_demo.ipynb`, both opening directly in Google Colab from the public GitHub companion. They walk through bundle loading, v1 checkpoint loading and inference, track-formatted submission writing, and local scoring. Release waves: starter kit + reference checkpoints + scorer at acceptance; hidden rows stay server-side at the scorer (never released).

1.7 Website, tutorial and documentation

Canonical entry points: the **public GitHub companion** `Reza2kn/uncgpt-2026-neurips` (track READMEs, starter kit, scorer, baseline notebooks, reproduction scripts, semantic non-collapse charts); the **public Hugging Face collection** aggregating datasets, the reference checkpoint, and the demo Space on one page (see Appendix references); and the **generation infrastructure mirror** `Reza2kn/KakoVerse`. A dedicated domain at `uncgpt.org` (reserved) is provisioned within two weeks of acceptance per Competition Track requirements.

A FAQ / Tutorial section ships at launch under `uncgpt.org/faq`, paired with the WebGPU browser demo Space so participants can probe the reference baseline interactively before submitting. Dedicated contact: `competitions@uncgpt.org` (mailbox provisioned at acceptance; reachable meanwhile at the organizer’s listed email). Pipeline architecture and rationale: `in-repo white papers proposal/draft_v13.md` and `proposal/pipeline_evidence_v13.md`.

2 Organizational aspects

2.1 Protocol

Joining. Participants register through the competition website using either a Hugging Face or GitHub identity (no third-party platform account required). Registration captures team name, primary contact email, and track(s) of interest. Data access requires no gating: bundle parquets are direct downloads from the public Hugging Face repos in Section 1.3.

Submission format. Participants submit either (a) a single JSONL prediction file conforming to the per-track schema described in `competition/scorer/README.md`, or (b) for the Dataset track, a parquet dataset bundle plus a YAML recipe overlay against the frozen retrain config, or (c) for the Federated track, a per-client weight-delta tarball plus a DP accountant report. Submissions are uploaded to a CodaBench instance configured against the same scorer used locally.

Evaluation. The CodaBench scorer is the exact `competition/scorer/score_submission.py` module from the public repo, pinned to a release tag. It (1) validates input schema, (2) runs scoring against hidden rows for tracks 1–4, (3) triggers the retrain runner against frozen seeds for track 5, (4) validates DP-accounting fields and applies the per-utility formula for track 6, (5) writes a public scorecard with bootstrap CIs.

Phases and platform. A single submission phase with a public leaderboard that displays final scores only after the submission window closes (to limit leaderboard-overfitting). During the window, participants see only their own validated score on the public bundles and a coarse ranking band (decile) against hidden rows.

Anti-cheating / anti-overfitting controls. (i) Hidden rows are never exposed; only aggregate scorecards return. (ii) Per-team submission cap (10 graded submissions during the window; unlimited dry-runs on public bundles only). (iii) Reproduction audit on top-5 finishers per track: organizers re-execute submitted code against frozen seeds on the standardized environment (`proposal/standardized_environment_v11.md`) and disqualify mismatches above a published tolerance. (iv) Plagiarism check across submissions to detect shared private wrappers.

Beta tests and dry runs. Two pre-launch dry runs are scheduled: (a) a scorer-only dry run on synthetic submissions in the two weeks after acceptance (validates CodaBench wiring); (b) a closed beta with 5–8 invited friendly teams in the two weeks before launch (validates documentation and starter notebooks against real participant friction). Both dry runs use a separate evaluation cohort drawn from the same generator with different seeds, never the actual hidden rows.

Fair-play / level-playing-field provisions. A standardized reproduction environment (single A100 80GB, PyTorch 2.5, CUDA 13, Python 3.12) is published. Participants without A100 access can request **free reproduction passes** from the organizer compute pool (Section 3.2) to ensure resource asymmetry does not distort rankings.

2.2 Rules and Engagement

Draft contestant rules (verbatim).

1. By submitting, you agree to release your submission’s scoring artifacts (predictions, code, and metadata) under CC-BY-4.0 and your code under Apache-2.0 for organizer use in the post-competition workshop, results paper, and reproducibility audit.
2. All evaluation data, including hidden rows, remain organizer-owned. You may not attempt to recover hidden rows from the scorer; doing so disqualifies your team across all tracks.
3. Maximum 10 graded submissions per team per track during the submission window. Unlimited local dry-runs on public bundles are encouraged.
4. Pretraining and fine-tuning may use any data that is (a) publicly available before the official competition launch date and (b) either openly licensed or used under an explicit exception. Using leaked hidden rows is disqualifying.
5. Inference budget per submission: max 1,024 generated tokens per turn, max 64 turns per conversation, max 30 minutes wall-clock on the standardized environment. The Federated track adds a per-round compute cap published in the track README.
6. The Empathic / Disability-Care track requires explicit acknowledgement of the safety rubric and harm-trigger taxonomy. Submissions producing diagnoses, professional-authority imitation, coercive dependency, or premature hotline deflection are penalized per the published rubric; egregious or repeated violations are disqualifying.
7. Inclusivity: partial submissions (e.g., only a subset of eligible languages on Low-Resource; only a subset of clients on Federated) are accepted, scored on the submitted subset, and listed without penalty in the final leaderboard alongside complete submissions. We do not call out or shame participants who could not complete every track.
8. Code of Conduct: the NeurIPS Code of Conduct applies to all communication channels.

Why these rules. 1–3 protect hidden-cohort integrity and limit overfitting; 4 prevents trivial wins from leaked answers while keeping the open ecosystem in; 5 enforces constrained-compute realism; 6 keeps the empathic track from drifting into hotline-deflection theater or clinical pretense; 7 keeps the competition accessible to small teams; 8 imports NeurIPS’s standing norms.

Communication channels. Email competitions@uncgpt.org; GitHub Discussions on the companion repo; bi-weekly recorded office hours. Rule clarifications and deadline updates ship through a pinned Discussions thread, the mailing list, and dataset/model card revisions on Hugging Face. All updates are timestamped and never retroactive after a phase begins.

2.3 Schedule and readiness

Readiness at proposal time. The following are ready now and reviewable in the public repo / Hugging Face: persona pool (1M+), names registry (15k, zero gaps), live-approved 4,761-conversation gold pool, scorer source and tests, public bundles for all six tracks, reference scorecards, smoke checkpoint, and WebGPU demo Space. The UncGPT v1 full-pass training is currently in flight (proposal-freeze snapshot: $\approx 10k$ / 12,399 steps); the completed weights will be mirrored before launch.

Detailed post-acceptance timeline.

- **2026-05-15 AoE:** proposal freeze; all listed public assets and bundles tagged at frozen revisions.
- **2026-06-15:** official acceptance notification window (per NeurIPS Competition Track schedule).
- **2026-06-16 to 2026-06-30:** launch phase — public release note, scorer / asset consistency check, competition website live, CodaBench instance configured, scorer-only dry run executed.
- **2026-07-01 to 2026-09-30:** full submission window; bi-weekly office hours; closed beta with friendly teams in the first two weeks.
- **2026-10-01 to 2026-10-31:** final submission freeze; reproduction audits on top-5 per track; manual empathic review panel convenes; rankings finalized.
- **2026-11-01 to 2026-11-30:** results paper drafting; workshop talk slate confirmed.
- **2026-12-11 to 2026-12-12:** in-person NeurIPS Competition Track workshop presentation.

Contingency plan. If a track receives fewer than 5 distinct teams by 2026-08-15, organizers run a targeted mid-window outreach push; if engagement remains below threshold by 2026-09-15, that track folds into the workshop as an organizer-led baseline study (public bundles intact) rather than publishing a degenerate leaderboard. If CodaBench has unrecoverable downtime, a backup self-hosted scorer (already exercised in the cold-run audit) accepts submissions through a GitHub Actions workflow. Timeline slips greater than 7 days are communicated immediately on all channels with a revised schedule.

2.4 Competition promotion and incentives

Promotion plan. The call goes through the NeurIPS Competition Track call, the Hugging Face Datasets & Benchmarks newsletter, the AfricaNLP / Masakhane mailing lists (Low-Resource), Lat-inX in AI / Queer in AI / Women in ML / Black in AI channels (under-represented groups), SIG-DIAL and ACL safety-and-ethics communities (Empathic), the federated-learning research community (FLOWER, FedML — Federated), and invited talks at HuggingFace community calls, the EleutherAI research stage, and the Cohere C4AI seminar.

Incentives. **Recognition:** podium slot at the workshop for per-track winners and runners-up; named-author credit on the joint results paper (consortium author for any submitting+attending team; named author for top-3 per track who contribute a method writeup). **Compute:** free A100 reproduction passes from the organizer pool, capped per team. **Travel:** a small organizer-funded travel-grant pool for at least one member of each top-3 team that would otherwise be unable to attend, prioritized for participants from under-represented regions. **Monetary:** no mandatory cash prize at proposal time; any sponsor-committed prize at launch is disclosed transparently and does not affect ranking formulas.

Attracting under-represented groups. Low-Resource and Multilingual tracks lower the bar for groups historically excluded by English-only benchmarks. The Empathic / Disability-Care track is co-designed with disability-support reviewers (pre-launch panel). All starter notebooks run on the free Colab tier. Partial submissions are scored without penalty. Outreach uses community-led mailing lists, and we explicitly avoid scoring formulas that disadvantage teams without proprietary infrastructure.

2.5 Competition track workshop and dissemination

The Competition Track workshop slot will be used to: (i) present per-track winners (10-minute talks each, podium-only — no parallel sessions to avoid fragmenting attention); (ii) run a 30-minute organizer-led analysis of cross-track patterns (which interventions transferred, which collapsed); (iii) run a moderated panel with one community reviewer from the empathic pre-launch panel and one cross-track top-finisher to discuss what sitting with looked like in practice. A live Q&A board collects participant questions in the days before the workshop and during.

Dissemination and publication. (a) A joint organizer+-participant **results paper** will be drafted in November and submitted to a Datasets & Benchmarks venue (NeurIPS D&B for the following cycle, or TMLR for rolling acceptance), with authorship levels as described in Section 2.4. (b) All non-disqualified submissions are archived on Hugging Face with attribution. (c) A retrospective post on the competition website and Hugging Face blog summarizes lessons learned, surprising failure modes, and proposed future directions. (d) The hidden evaluation cohort is not released post-competition — it remains the canonical hidden eval for a possible 2027 follow-up.

3 Resources

3.1 Organizing team

The lead organizer (Reza Sayar) is the primary accountable owner for submission, the release contract, scoring, notebooks, and audit evidence; their biography is in Appendix A. The six-track scope is feasible because all tracks share a single infrastructure stack and one manifest-controlled asset family — this is not six independent competitions, it is one competition with six controlled perturbations on the same substrate.

The proposal is submitted under a single confirmed organizer identity at proposal time. External pre-launch lanes are bounded and named:

- **Empathic / Disability-Care pre-launch panel** (3–5 community reviewers, including disability-care practitioners and at least one crisis-line-experienced reviewer). Recruitment is in-progress through community channels.
- **Multilingual / Low-Resource per-language reviewers** (1–2 native speakers per eligible language for register-fidelity audit on a sampled subset of hidden rows).
- **Federated protocol sanity reviewer** (1 federated-learning practitioner to validate the DP-accounting contract before launch).

The separation between confirmed execution surface (lead organizer) and bounded pre-launch review lanes (named external reviewers) is documented in `proposal/committee_matrix_v12.md`.

3.2 Resources provided by organizers

Compute: a small dedicated NVIDIA A100 80GB pool for reference-baseline training/inference, scorer execution, dry runs, reproduction audits, and free reproduction passes for participants without A100 access; plus a CPU + commodity-GPU pool for CodaBench scoring. **Hosting:** Hugging Face (datasets, models, Space); GitHub (code, docs, notebooks); a dedicated competition site (`uncgpt.org`, reserved); CodaBench (submission, leaderboard). **Software:** the full generation harness, gate stack, scorer, baselines, notebooks, and reproduction scripts are organizer-built and openly licensed. **Staff:** lead organizer plus named pre-launch reviewers (Section 3.1); recruitment of additional support for the submission-window office-hours rotation is in-progress.

3.3 Support requested

CodaBench hosting: a NeurIPS-sponsored CodaBench instance for the submission window, or equivalent platform credit. **In-person workshop slot:** the standard NeurIPS Competition Track in-person slot (we acknowledge that the Competition Track is an in-person event) with A/V for the 10-minute per-track winner talks and the moderated panel. **Travel-grant coordination:** coordination with the NeurIPS travel-grant or equivalent program to backstop the organizer-funded travel pool described in Section 2.4, prioritized for participants from under-represented regions. **Cross-promotion:** inclusion in the Competition Track call distribution, NeurIPS social-media announcements, and the Datasets & Benchmarks newsletter.

A Biography of all team members

Reza Sayar (Lead Organizer). Independent researcher and engineer working on small, runnable, multilingual conversational models for caregiving-adjacent settings. Prior open-source releases on

Hugging Face under huggingface.co/Reza2kn include the Persian RaahNaameh-1 textual corpus, the OLDI Wikipedia MT seeds for Persian and English–Persian, the Common Voice 24 Persian release, the XRAY model line, multiple NVFP4 / AWQ quantized Persian-speech model releases, and the KakoVerse persona-and-conversation-generation substrate. Direct lived experience supporting people through crisis informs the empathic-stance design of this competition — the explicit goal is to evaluate small models that can sit with a person rather than deflect, while still respecting bounded safety and never imitating clinical authority. OpenReview profile: [~Reza_Sayar1](https://openreview.net/profile/~Reza_Sayar1). GitHub: github.com/Reza2kn.

Pre-launch external reviewers (recruitment in-progress; bounded scope, not proposal-time confirmed execution surface). Empathic / Disability-Care reviewers (3–5, including disability-care practitioners and at least one reviewer with crisis-support practice); Multilingual / Low-Resource per-language reviewers (1–2 native speakers per eligible language for register-fidelity audit on sampled hidden rows); Federated protocol sanity reviewer (1 federated-learning practitioner for DP-accounting contract validation). Final biographies will be added to the website and to the results paper at launch and post-competition respectively, with explicit consent.

References

- [1] NeurIPS 2026 Call for Competitions. <https://neurips.cc/Conferences/2026/CallForCompetitions>
- [2] NeurIPS Code of Ethics. <https://neurips.cc/public/EthicsGuidelines>
- [3] Sayar, Reza. OpenReview profile. https://openreview.net/profile?id=~Reza_Sayar1
- [4] Sayar, Reza. Hugging Face organization page. <https://huggingface.co/Reza2kn>
- [5] UncGPT NeurIPS Competition. Hugging Face collection (datasets, reference checkpoints, demo Space): huggingface.co/collections/Reza2kn/uncgpt-2026-neurips-competition.
- [6] Reza2kn/uncgpt-2026-neurips. GitHub companion repository to this proposal (track READMEs, scorer, baselines, notebooks, reproduction scripts). <https://github.com/Reza2kn/uncgpt-2026-neurips>
- [7] Reza2kn/uncgpt-personas. Hugging Face dataset (1,000,069 personas; 15,000 names; gold pools; semantic non-collapse summary). <https://huggingface.co/datasets/Reza2kn/uncgpt-personas>
- [8] Reza2kn/uncgpt-conversations-v7-69-total. Hugging Face dataset (4,761 live-approved conversations, 11 languages, 69 skills). <https://huggingface.co/datasets/Reza2kn/uncgpt-conversations-v7-69-total>
- [9] Reza2kn/uncgpt-conversations-semantic-approved-1p25. <https://huggingface.co/datasets/Reza2kn/uncgpt-conversations-semantic-approved-1p25>
- [10] Reza2kn/uncgpt-conversations-semantic-approved-1p25-repaired-paperclip. <https://huggingface.co/datasets/Reza2kn/uncgpt-conversations-semantic-approved-1p25-repaired-paperclip>
- [11] Reza2kn/uncgpt-conversations-semantic-approved-1p50-candidate. <https://huggingface.co/datasets/Reza2kn/uncgpt-conversations-semantic-approved-1p50-candidate>
- [12] Reza2kn/uncgpt-conversations-informal-approved-1p25. <https://huggingface.co/datasets/Reza2kn/uncgpt-conversations-informal-approved-1p25>
- [13] Reza2kn/uncgpt-persian-ultraclean-review-2026-05-14. <https://huggingface.co/datasets/Reza2kn/uncgpt-persian-ultraclean-review-2026-05-14>
- [14] Reza2kn/uncgpt-69-smoke-checkpoints-20260515. Hugging Face model (UncGPT v1 step-100 smoke checkpoint, tokenizer, configs, source). <https://huggingface.co/Reza2kn/uncgpt-69-smoke-checkpoints-20260515>
- [15] Reza2kn/uncgpt-69-hybrid-demo. Hugging Face Space (in-browser WebGPU inference demo). <https://huggingface.co/spaces/Reza2kn/uncgpt-69-hybrid-demo>
- [16] Reza2kn/KakoVerse. GitHub repository (RL environments and persona / conversation generation substrate). <https://github.com/Reza2kn/KakoVerse>

- [17] Sayar, Reza. GitHub profile. <https://github.com/Reza2kn>
- [18] Reza2kn/RaahNaameh-1-textual-corpus. Hugging Face dataset. <https://huggingface.co/datasets/Reza2kn/RaahNaameh-1-textual-corpus>
- [19] Reza2kn/OLDI-Wikipedia-MTSeed-Persian. Hugging Face dataset. <https://huggingface.co/datasets/Reza2kn/OLDI-Wikipedia-MTSeed-Persian>
- [20] Reza2kn/OLDI-Wikipedia-MTSeed-English-Persian. Hugging Face dataset. <https://huggingface.co/datasets/Reza2kn/OLDI-Wikipedia-MTSeed-English-Persian>
- [21] Reza2kn/commonvoice-24-fa. Hugging Face dataset. <https://huggingface.co/datasets/Reza2kn/commonvoice-24-fa>